

Creative Narrative Transformation Systems (CNTS): A Layered, Orchestrated, and Human-in-the-Loop Framework for AI-Driven Multimedia Generation

Vijaya K Nadendla^{1,†} 

* Correspondence: vnadendla@gmail.com; Tel.: +01-6095104995

† Current address: 51 Linwood Dr Monroe Township NJ USA 08831.

Abstract

Recent advances in generative AI enable direct synthesis of multimedia content from textual inputs. However, these systems lack structured mechanisms for translating narrative intent into coherent cinematic output. We introduce Creative Narrative Transformation Systems (CNTS), a framework grounded in the Foundational Reasoning and Feedback Protocol (FRFP), which separates explicit artifacts from tacit human judgment. CNTS combines structured narrative transformation, layered cinematic decomposition, and multi-primary action orchestration. Rather than producing single outputs, the system constructs structured spaces of candidate interpretations for human evaluation. This approach aligns creative generation with principles of epistemic governance, enabling controllable, interpretable, and collaborative multimedia production.

Keywords: human–AI collaboration; explicit–tacit boundary; endorsement; bad faith; observational equivalence; governance; credence inflation; institutional capture

0. Introduction

The rapid evolution of generative artificial intelligence has fundamentally altered the landscape of media creation. Systems capable of producing images, video, and audio from textual prompts have made it possible to bypass traditional production pipelines and directly synthesize cinematic artifacts. While these advances represent a significant technological milestone, they expose a deeper conceptual limitation: current systems lack a principled method for translating narrative intent into coherent, layered cinematic expression.

A director’s narrative typically encodes story elements such as characters, events, and emotional arcs. However, cinema is not merely a representation of narrative; it is an emergent form constructed through multiple interacting layers, including visual composition, color, lighting, sound design, temporal pacing, and actor interaction. These layers are rarely specified explicitly in narrative form and are instead realized through interpretive decisions made by directors, cinematographers, editors, and performers.

Contemporary AI systems attempt to bridge this gap through prompt-based generation. In such systems, textual instructions implicitly encode multiple creative decisions simultaneously. This approach, while flexible, leads to two fundamental problems. First, it collapses multiple valid interpretations into a single output without exposing the underlying assumptions. Second, it lacks mechanisms for coordinating the complex, multi-layered nature of cinematic scenes, particularly those involving multiple actors and simultaneous actions.

As a result, AI-generated multimedia often exhibits one or more of the following failure modes: visual incoherence, tonal inconsistency, shallow emotional impact, or unrealistic interaction between elements within a scene. These failures are not merely technical limitations but arise from a deeper absence of structured reasoning about creative ambiguity and layered representation.

This paper proposes that the problem of AI-driven multimedia generation is not solely a question of improving models, but of introducing a framework for structured creative reasoning. Specifically, we argue that AI systems must transition from generating single outputs to constructing *spaces of possible interpretations*, each explicitly defined, inspectable, and subject to human evaluation.

To address this, we introduce the **Creative Narrative Transformation System (CNTS)**, a framework that formalizes the transformation of narrative into cinematic artifacts through three key components: (i) a structured narrative transformation protocol, (ii) a layered model of cinematic representation, and (iii) an orchestration model for multi-actor scenes. The system operates under a human-in-the-loop paradigm in which AI generates candidate interpretations and human creators retain authority over selection and refinement.

The conceptual foundation of this work originates in the **Foundational Reasoning and Feedback Protocol (FRFP)**, which introduces a formal separation between explicit representations and tacit human judgment. While originally developed for analytical reasoning, its principles extend naturally to creative systems, where ambiguity and interpretation are central.

1. From FRFP to Creative Systems

The Creative Narrative Transformation System (CNTS) is grounded in the Foundational Reasoning and Feedback Protocol (FRFP), a framework originally developed to address failures in analytical reasoning systems. To understand CNTS, it is essential to first understand the core ideas of FRFP in intuitive terms.

FRFP begins with a simple but powerful observation: not all forms of knowledge can be treated the same way. Some knowledge can be written down, transmitted, and processed by machines. Other forms of knowledge—such as judgment, interpretation, and understanding—exist only within human cognition.

FRFP formalizes this distinction as:

- **Explicit artifacts:** representations that can be written, stored, and processed (e.g., text, prompts, models, images)
- **Tacit states:** internal human capacities such as understanding, judgment, and contextual awareness

For example, a movie script is an explicit artifact. However, deciding whether a scene is emotionally powerful or culturally appropriate is a tacit judgment.

FRFP enforces a strict rule: *machine systems must operate only on explicit artifacts, while final judgment must remain in the tacit domain*. This prevents systems from implicitly claiming correctness or authority beyond what can be justified.

A key result of FRFP is that any system relying only on explicit representations will eventually fail over long reasoning chains. In analytical contexts, this appears as incorrect conclusions. In creative contexts, it appears as outputs that look plausible but feel “wrong”—for example, a scene that is visually impressive but emotionally empty.

FRFP also identifies a common failure pattern in modern AI workflows: the collapse of multiple steps into a single output. In typical prompt-based systems, generation, interpretation, and endorsement are fused into one response. This creates the illusion that the system “knows” what it is doing, when in reality it is producing one possible interpretation without exposing alternatives.

FRFP addresses this through two principles:

1. **Branching:** multiple interpretations must be generated and preserved
2. **Exposure:** assumptions must be made visible rather than hidden

Applying this to creative systems leads to a fundamental shift:

Instead of asking AI to generate the best scene, we ask it to generate multiple plausible scenes, each representing a different interpretation.

FRFP also defines a structured interaction loop between human and machine:

- Representation Initiation (RI): structuring the input
- Explicit Computation (EC): generating outputs
- Representational Backflow (RB): presenting outputs
- Tacit Evaluation (TE): human judgment
- Human Final Decision (HFD): selection and endorsement

In CNTS, this loop directly maps to the creative process. The AI generates multiple cinematic possibilities, and the human evaluates and selects among them.

Thus, CNTS is not a separate idea but a direct extension of FRFP into the domain of multimedia generation.

2. The Nature of Cinematic Representation

To understand why AI struggles with video generation, it is necessary to understand how cinema differs from narrative.

Narrative describes events in a sequential and abstract way. For example, a narrative might state: “a woman hesitates before opening a letter.” This description specifies what happens, but not how it is experienced.

Cinema, by contrast, is not just about events. It is about how those events are presented to the viewer. This presentation is constructed through multiple layers that operate simultaneously.

For clarity, we can think of these layers as follows:

- **Visual composition:** where things are placed in the frame
- **Lighting and color:** how the scene is illuminated and what mood it conveys
- **Sound:** music, silence, and ambient noise
- **Timing:** how quickly or slowly events unfold
- **Actor behavior:** gestures, expressions, and interactions

Consider the same narrative action—“a woman opens a letter.” This can be rendered in many ways:

- A quiet close-up with no music (intimate, reflective)
- A fast cut with rising sound (tense, suspenseful)
- A wide shot with dramatic lighting (theatrical, symbolic)

Each version tells the same story, but creates a different emotional experience.

This leads to an important conclusion:

There is no single correct cinematic representation of a narrative.

Instead, there exists a space of valid interpretations.

Current AI systems struggle because they attempt to compress all these layers into a single prompt. This causes uncontrolled interactions between elements—for example, mismatched lighting and emotion, or inconsistent actor behavior.

To address this, CNTS introduces the Layered Creative Collaboration Model (LCCM), which separates these layers and allows them to be controlled independently.

A second challenge arises in scenes with multiple actors. In real scenes, multiple things happen at the same time. For example:

- one character performs an action
- another observes
- a third reacts emotionally

These are not secondary details—they are essential to the scene.

This motivates the Multi-Primary Orchestration Model (MPOM), which treats a scene as a system of concurrent actions rather than a single event.

2.1. Intuition Behind MPOM

The MPOM model may appear mathematical, but its core idea is simple.

In real scenes, multiple important actions can happen at the same time. For example:

- a woman lifts a lid
- a man watches closely
- another person reacts emotionally

Traditional systems force one of these to be “primary.” MPOM instead allows multiple actions to be important simultaneously.

Each action is assigned:

- a weight (how important it is)
- a timing (when it happens)

The system then ensures that:

- the scene remains understandable
- attention is guided (through camera, lighting, etc.)

This makes generated scenes feel more natural and less artificial. Together, LCCM and MPOM allow CNTS to represent cinema as a structured system rather than a single instruction.

3. Creative Narrative Transformation System Architecture

Having established the epistemic foundation and representational structure, we now formalize the architecture of the Creative Narrative Transformation System (CNTS). The system is designed as a layered transformation pipeline operating entirely within the explicit artifact space, while deferring interpretation and endorsement to the human domain, in accordance with FRFP principles.

3.1. System Overview

The CNTS system can be understood as a pipeline that gradually transforms a narrative into video outputs.

Narrative → Structure → Layers → Actions → Prompts → Video → Selection

Each stage performs a specific role:

- **Narrative Structuring:** converts story into scenes and key moments
- **Layer Decomposition:** separates visual, sound, and emotional aspects
- **Action Orchestration:** models how multiple actors behave
- **Prompt Composition:** translates structured data into prompts
- **Generation:** produces video outputs
- **Selection:** human chooses best results

Importantly, the system does not attempt to produce a final result in one step. Instead, it produces multiple candidates and relies on human judgment to select among them. To make this concrete, consider a simple example. A narrative describing a tense dinner scene is first broken into key moments (arrival, conversation, reveal). Each moment is then expanded into visual and emotional layers (lighting, sound, pacing). These layers are combined with action orchestration to produce structured prompts, which are used to generate multiple video variations. The human creator then selects and refines the best outputs.

3.2. Mapping to FRFP Interaction Primitives

The CNTS pipeline can be expressed as an instantiation of FRFP interaction primitives:

- **Representation Initiation (RI):** Narrative input is transformed into structured elements such as scenes, characters, and relationships.

- **Explicit Computation (EC):** The system generates layered representations, action graphs, and prompts.
- **Representational Backflow (RB):** Candidate outputs and intermediate artifacts are presented for inspection.
- **Tacit Evaluation (TE):** Human evaluates coherence, emotional impact, and cultural alignment.
- **Human Final Decision (HFD):** Human selects or refines the final artifact.

This mapping ensures that CNTS adheres to the explicit–tacit separation enforced by FRFP, preventing implicit endorsement within machine-generated outputs.

3.3. Component Definitions

Narrative Structuring (CNTP)

This component extracts structured representations from narrative input, including:

- key scenes
- emotional arcs
- character relationships
- narrative constraints (e.g., no spoilers in trailers)

The output is a set of explicit narrative artifacts suitable for further transformation.

Layer Decomposition (LCCM)

The system decomposes each scene into independent layers:

- visual composition
- color and lighting
- sound and music
- temporal pacing
- emotional modulation

Each layer is represented explicitly, enabling controlled variation and preventing implicit coupling.

Action Orchestration (MPOM)

Scenes are represented as networks of concurrent actions. Each action is modeled as a node with associated properties, enabling coordination across multiple actors.

Prompt Composition

Structured prompts are generated by combining layer specifications and action graphs. Prompts remain interpretable and reproducible.

3.4. Why Prompt Structure Matters

In many AI systems, prompts are treated as free-form text. However, this leads to inconsistent results because multiple decisions are mixed together.

CNTS treats prompts as structured representations.

Instead of writing:

“a dramatic scene with tension”

we break it into components:

- what is happening (actions)
- how it looks (visual style)
- how it moves (camera)
- how it sounds (audio)

This makes prompts:

- easier to control

- easier to modify
- easier to reproduce

For example, we can change only lighting without changing actions, which is not possible with unstructured prompts.

Generation and Iteration

The system produces multiple candidate outputs using generative models. Variations are introduced through parameter changes, prompt modifications, and seed variation.

3.5. Iteration and Selection Loop

The CNTS system operates iteratively:

Generate → Compare → Select → Refine → Repeat

This loop ensures that quality emerges through comparison and refinement rather than single-shot generation.

3.6. Epistemic Guarantees

By construction, CNTS satisfies the following properties:

- No implicit endorsement within explicit artifacts
- Preservation of ambiguity through branching
- Inspectability of intermediate representations
- Human authority over final decisions

These properties follow directly from FRFP constraints on explicit–tacit separation.

3.7. From Intuition to Formal Model

We now formalize the intuitive description of multi-actor scenes introduced earlier. The goal is not to introduce abstract mathematics, but to provide a precise representation of the same idea: that multiple actions can occur simultaneously and contribute jointly to the meaning of a scene.

Each element of the formal model corresponds directly to observable components in a cinematic scene.

4. Multi-Primary Orchestration Model (MPOM)

We now formalize the Multi-Primary Orchestration Model (MPOM), which addresses the representation of multi-actor cinematic scenes.

4.1. Motivation

Traditional prompt-based systems implicitly assume a single dominant action per scene. While this simplifies generation, it fails to capture the complexity of real cinematic interactions, where multiple actors perform concurrent actions that contribute jointly to meaning.

4.2. Action Graph Representation

We define a scene as a directed graph:

$$G = (V, E)$$

where:

- V is a set of action nodes
- E is a set of relational edges

Each node $v \in V$ is defined as:

$$v = (a, \alpha, w, t)$$

where:

- a = actor
- α = action
- $w \in \mathbb{R}^+ = \text{intensity (weight)}$
- $t \in \mathbb{R} = \text{temporal offset}$

4.3. Worked Example

To illustrate how the formal model corresponds to a real scene, consider a simple multi-actor interaction:

- a widow lifts a lid
- a host observes her closely
- a daughter reacts emotionally

This scene can be represented as:

$$v_1 = (\text{Widow}, \text{lift lid}, 0.9, 0.0)$$

$$v_2 = (\text{Host}, \text{observe}, 0.8, 0.2)$$

$$v_3 = (\text{Daughter}, \text{react}, 0.6, 0.5)$$

In this configuration, multiple actions exceed the threshold θ , making them jointly primary.

The temporal offsets introduce natural sequencing, while the attention function distributes viewer focus across these actions. Cinematic techniques such as framing and lighting determine the specific values of $A(v)$ in practice.

4.4. Multi-Primary Condition

We define a subset $P \subseteq V$ as the set of primary actions:

$$P = \{v \in V \mid w(v) \geq \theta\}$$

for some threshold θ .

Unlike hierarchical models, $|P| \geq 1$, allowing multiple primary actions.

4.5. Temporal Coordination

Temporal relationships are encoded through node offsets:

$$t : V \rightarrow \mathbb{R}$$

These offsets introduce micro-delays that produce naturalistic sequencing of actions.

4.6. Attention Function

To maintain interpretability, we define an attention function:

$$A : V \rightarrow [0, 1]$$

such that:

$$\sum_{v \in V} A(v) = 1$$

This function represents the distribution of viewer attention across actions. Attention is influenced by cinematic parameters such as framing, lighting, and motion.

4.7. *Orchestration Constraint*

A valid scene must satisfy:

$$\max_{v \in V} A(v) \geq \gamma$$

for some γ , ensuring that at least one action remains clearly interpretable.

4.8. *Emergent Meaning*

The meaning of a scene is not determined by any single node but emerges from the combined configuration:

$$M = f(G, A)$$

where f represents the viewer's interpretive process, which depends on perceptual, cultural, and emotional factors.

4.9. *Integration with LCCM*

Each node in the action graph is associated with layer-specific parameters:

- visual emphasis (framing, focus)
- lighting intensity
- sound cues
- timing adjustments

Thus, MPOM operates jointly with LCCM to produce coherent cinematic outputs.

4.10. *Discussion*

MPOM generalizes single-action models by allowing multiple concurrent focal points while maintaining interpretability through attention constraints. This aligns with real-world cinematic practice, where meaning emerges from coordinated interactions rather than isolated actions.

From an FRFP perspective, MPOM ensures that multi-actor interactions remain explicit and inspectable, preventing implicit collapse of complex dynamics into a single representation.

5. **Iterative Generation and Prompt Semantics**

The preceding sections establish the structural and representational foundations of CNTS. We now turn to the operational layer: how structured representations are translated into generated media through prompts and iterative workflows.

5.1. *Prompt Semantics as Representation Encoding*

In CNTS, prompts are not treated as ad hoc instructions but as structured encodings of explicit artifacts. A prompt is a composition of multiple representational components derived from CNTP, LCCM, and MPOM.

We define a prompt P as:

$$P = (S, V, L, C, M, A)$$

where:

- S = shot type (composition)
- V = action set (from MPOM graph)
- L = layer configuration (LCCM)
- C = camera parameters
- M = motion or temporal dynamics
- A = audio specification

This structure ensures that prompts remain interpretable, decomposable, and reproducible.

5.2. Prompt Decomposition

Each component of P corresponds to a specific layer:

- $S \rightarrow$ visual composition layer
- $V \rightarrow$ action orchestration layer
- $L \rightarrow$ color, lighting, and style
- $C \rightarrow$ camera control
- $M \rightarrow$ temporal modulation
- $A \rightarrow$ sound layer

This mapping prevents implicit coupling between layers and allows controlled variation.

5.3. Iteration as Core Mechanism

Unlike traditional generation systems that aim for a single output, CNTS adopts an iterative generation model.

Let $\mathcal{P}_0$ be an initial prompt. The system generates a set of variations:

$$\{P_1, P_2, \dots, P_n\}$$

where each P_i is derived from $\mathcal{P}_0$ through controlled perturbations.

Variations may include:

- parameter adjustments
- layer modifications
- seed variation
- action weight changes

Each prompt produces a corresponding output:

$$O_i = G(P_i)$$

where G is the generative model. For example, an initial prompt describing a dinner scene may be varied by changing lighting (warm vs cold), camera movement (static vs slow push-in), or timing (faster vs slower pacing). Each variation produces a different emotional effect.

5.4. Selection and Refinement

Outputs are evaluated through a human-in-the-loop process:

$$O^* = \arg \max_{O_i} E(O_i)$$

where E represents tacit evaluation by a human.

The selected output O^* is then used to generate a refined prompt:

$$P' = \Phi(P^*, O^*)$$

where Φ is a refinement operator.

5.5. Connection to FRFP

This iterative process corresponds directly to FRFP primitives:

- EC: generation of candidate outputs
- RB: presentation of outputs
- TE: human evaluation
- HFD: selection and refinement

Importantly, evaluation remains outside the explicit system, preserving the explicit–tacit boundary.

5.6. *Emergent Quality Through Iteration*

Quality is not assumed to be present in any single output. Instead, it emerges through:

- comparison across variations
- selection of high-performing candidates
- iterative refinement

This aligns with practitioner observations that volume and structured variation outperform single-shot generation.

6. Evaluation, Failure Modes, and Epistemic Limits

We now examine the evaluation of CNTS outputs and the inherent limitations imposed by the explicit–tacit separation.

6.1. *Evaluation Framework*

Evaluation operates across multiple dimensions:

6.1.1. Structural Evaluation

- narrative clarity
- consistency across scenes

6.1.2. Cinematic Evaluation

- visual coherence
- pacing and timing

6.1.3. Emotional Evaluation

- engagement
- tension and release

6.1.4. Human Validation

- cultural authenticity
- emotional correctness

The final category necessarily resides in the tacit domain and cannot be automated.

6.2. *Failure Modes*

We identify several common failure modes:

Ambiguity Collapse

Premature selection of a single interpretation, leading to loss of creative diversity.

Layer Coupling

Implicit interaction between layers causing inconsistent or conflicting outputs.

Orchestration Failure

Incoherent multi-actor interactions due to lack of structured coordination.

Emotional Flattening

Outputs that are visually correct but lack emotional impact, often due to missing or poorly specified audio layers.

Overfitting to Prompts

Excessive reliance on prompt templates leading to generic outputs.

6.3. Mitigation Strategies

CNTS mitigates these failures through:

- enforced branching
- explicit layer decomposition
- action graph representation
- iterative generation and selection

6.4. Epistemic Limits from FRFP

FRFP establishes that no explicit-only system can guarantee correctness across all long-horizon interactions. This result extends to creative systems.

In particular:

- No explicit metric can fully evaluate emotional or aesthetic quality.
- No automated system can guarantee alignment with narrative intent in all cases.
- Some outputs will inevitably cross the explicit–tacit boundary without sufficient grounding.

These limitations are not implementation flaws but structural properties of the system.

6.5. Implications

The implications are twofold:

1. Human involvement is not optional but structurally necessary.
2. Evaluation must remain partially external to the system.

6.6. Reframing Success

Given these constraints, success in CNTS is not defined as producing a single optimal output, but as constructing a structured space of candidate outputs that:

- preserve ambiguity
- expose assumptions
- enable human selection

6.7. Final Statement

The CNTS framework does not eliminate uncertainty or subjectivity. Instead, it provides a structured way to manage them, ensuring that creative systems remain both expressive and accountable within the constraints of human–AI collaboration.

7. Case Study: The Last Supper Invitation

To illustrate the CNTS pipeline end-to-end, we walk through the transformation of a single narrative into a trailer sequence. This example demonstrates how abstract narrative input is progressively structured, layered, and rendered into multiple cinematic interpretations. We present a case study based on the narrative concept *The Last Supper Invitation*, a psychological ensemble drama centered around a high-stakes dinner scenario.

7.1. Narrative Input

The narrative involves a wealthy patriarch inviting a group of individuals with shared history to a formal dinner. Each guest faces a decision involving hidden consequences. The narrative emphasizes tension, unresolved relationships, and moral confrontation.

7.2. CNTP Application

The narrative is transformed into structured elements:

- **Key scenes:** arrival, table reveal, rule explanation, widow exit, escalation
- **Emotional spine:** tension → uncertainty → disruption → confrontation
- **Constraints:** maintain mystery, avoid revealing outcomes

7.3. LCCM Layer Decomposition

Each scene is decomposed:

- **Visual:** long dining table, symmetry, empty chairs
- **Color:** warm candlelight vs deep shadows
- **Sound:** silence, ambient tension, gradual score build
- **Temporal:** slow opening, accelerating cuts
- **Emotion:** unease, tension, escalation

7.4. MPOM Action Graph

A representative scene (widow event) is modeled as:

$$G = \{v_1, v_2, v_3\}$$

- v_1 : widow lifts lid ($w = 0.9, t = 0.0$)
- v_2 : host observes ($w = 0.8, t = 0.2$)
- v_3 : daughter reacts ($w = 0.6, t = 0.5$)

This multi-primary configuration creates tension through concurrent actions.

7.5. Prompt Construction

A structured prompt derived from CNTS:

wide cinematic shot, an older woman lifting a silver cloche,
a man watching intensely, a younger woman reacting emotionally,
moody candlelight, slow push-in, silence with ambient sound

7.6. Iteration and Selection

Multiple variations are generated by modifying:

- lighting intensity
- camera motion
- timing offsets

Human evaluation selects outputs based on:

- clarity of tension
- emotional impact
- visual coherence

7.7. Outcome

The resulting trailer is not a single deterministic output but a curated composition of selected candidate clips. This demonstrates CNTS as a system for structured exploration rather than direct generation.

8. Future Work and Extensions

CNTS provides a foundational framework, but several extensions remain open for exploration.

8.1. Multi-Agent Creative Systems

Future systems may distribute CNTS components across specialized agents:

- **Narrative agent:** structure and interpret story
- **Visual agent:** generate imagery
- **Critic agent:** evaluate outputs
- **Orchestrator agent:** coordinate workflow

This aligns with emerging research in multi-agent creative systems.

8.2. User Interface Design

A practical implementation requires interfaces that expose:

- branching interpretations
- layer controls
- action graphs
- variation comparisons

Such interfaces would enable intuitive human control over complex generative systems.

8.3. Tool Integration

CNTS can be integrated with existing tools:

- video generation systems
- editing software
- asset libraries

The framework acts as a coordination layer rather than a replacement.

8.4. Formal Extensions

Future work may include:

- formal proofs of MPOM properties
- optimization of attention functions
- probabilistic modeling of viewer perception

8.5. Evaluation Research

A major open question is how to measure creative quality. Hybrid approaches combining structural metrics and human evaluation may be required.

9. Appendix

9.1. Example Prompt Template

SHOT TYPE +
ACTION SET +
STYLE +
CAMERA +
AUDIO

9.2. Example Action Graph

$$v_1 = (\text{Widow, lift lid}, 0.9, 0.0)$$

$$v_2 = (\text{Host, observe}, 0.8, 0.2)$$

$$v_3 = (\text{Daughter, react}, 0.6, 0.5)$$

9.3. System Diagram




Figure 1. CNTS Transformation Pipeline

9.4. Implementation Notes

- prompts should encode only one dominant transformation at a time
- multiple variations should be generated per scene
- human evaluation is required for final selection

9.5. FRFP Alignment Summary

CNTS maintains:

- explicit artifact generation (AI)
- tacit evaluation (human)
- separation between generation and endorsement

This ensures compatibility with FRFP constraints on human–AI collaboration.

References

1. V. Nadendla. *Foundational Reasoning and Feedback Protocol (FRFP): A Formal Architecture for Long-Horizon Human–AI Collaboration*. Preprint, 2025.
2. V. Nadendla. *FRFP Limits in Human–AI Collaboration: A Position Paper on Capability, Fragility, and Socioeconomic Consequences*. Preprint, 2025.
3. V. Nadendla. *FRFP White Papers: Beyond Prompt Engineering and Epistemic Governance*. Preprint, 2025.
4. M. Polanyi. *The Tacit Dimension*. Routledge & Kegan Paul, 1966.
5. S. Amershi, D. Weld, M. Vorvoreanu, A. Fourney, B. Nushi, P. Collisson, J. Suh, S. Iqbal, P. Bennett, K. Inkpen, J. Teevan, R. Kikin-Gil, and E. Horvitz. Guidelines for human–AI interaction. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, 2019.
6. J. Park, J. O’Brien, C. Cai, M. Morris, P. Liang, and M. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)*, 2023.
7. Anonymous. *CREA: A Multi-Agent Framework for Creative Generation*. arXiv preprint, 2025.
8. Y. Zhang et al. Multimodal generative models for video synthesis: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2024.
9. B. Shneiderman. Human-centered artificial intelligence: Reliable, safe and trustworthy. *International Journal of Human–Computer Interaction*, 2020.